

# Implementation and Evaluation of a Congestion Manager

*David G. Andersen, Deepak Bansal, Dorothy Curtis  
Srinivasan Seshan, Hari Balakrishnan*

**MIT Laboratory for Computer Science**

<http://nms.lcs.mit.edu/projects/cm/>

October 2000

# End-to-End Congestion Control

---

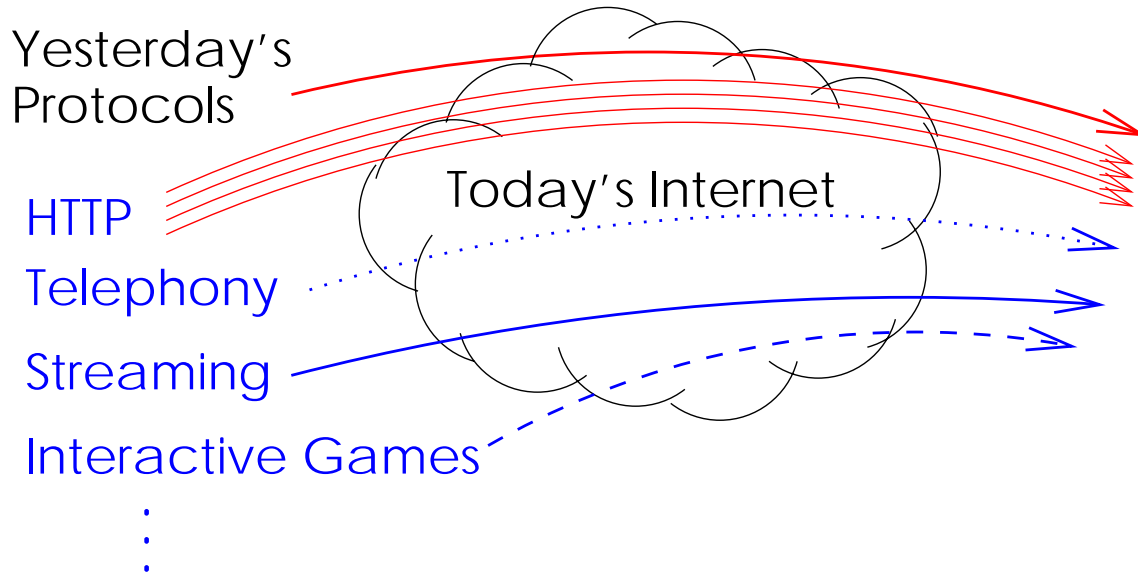
- ✓ Prevents congestion collapse
- ✓ Adapts to network conditions, other streams

Today's solution is TCP

- “AIMD” window-based congestion control
- Provides reliable, in-order bytestream

# Problems with current solution

---



When everything used TCP, it was sufficient

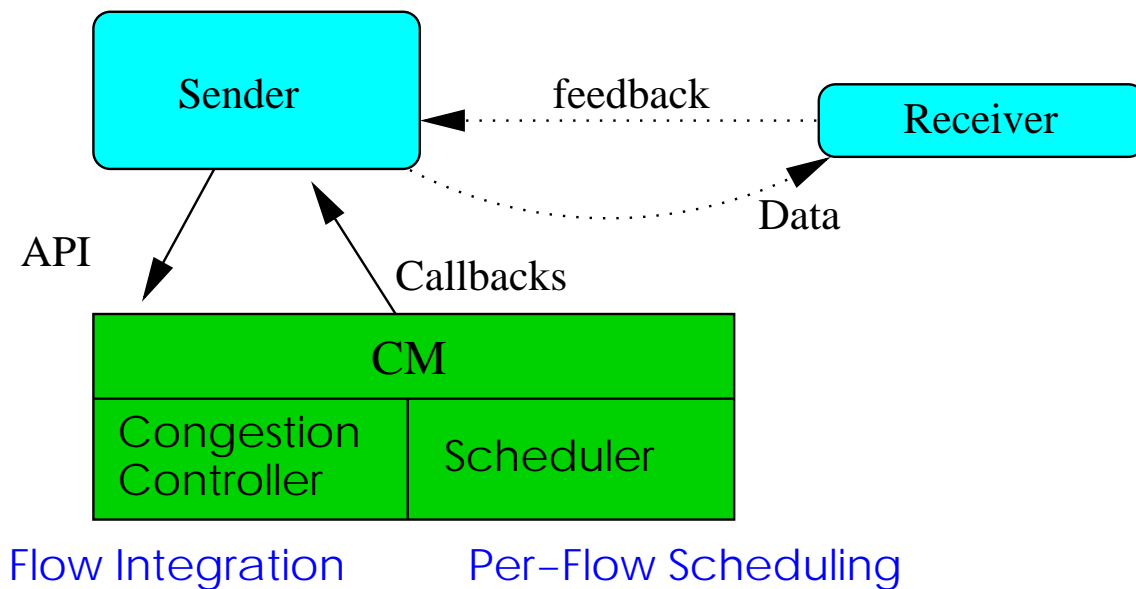
Today's traffic has moved beyond this:

- ✗ Not everything wants TCP
- ✗ Multiple TCP streams are less social

# The solution: The Congestion Manager

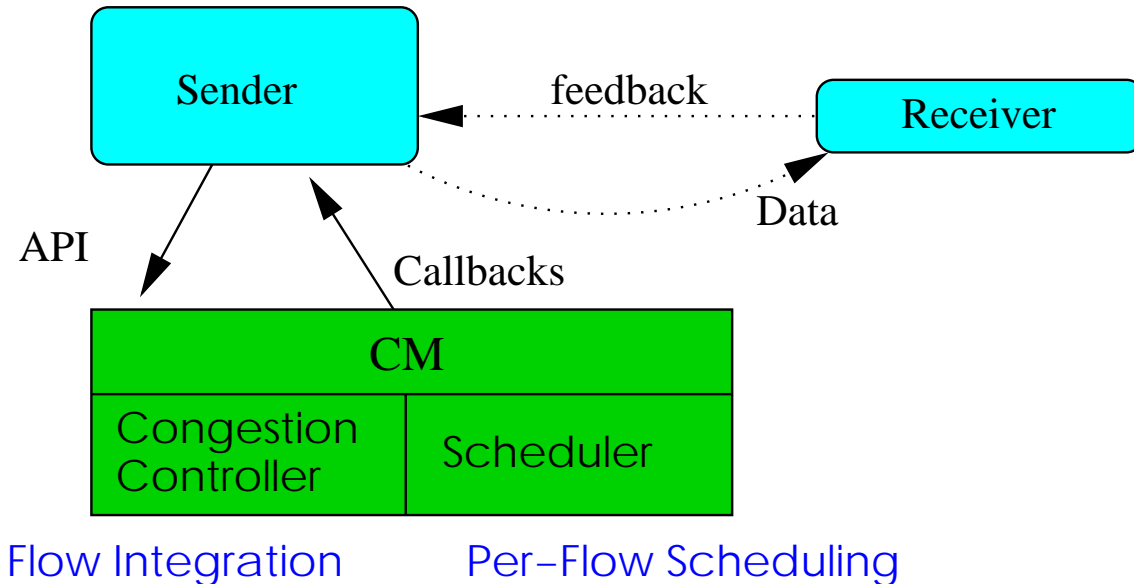
Sigcomm 99 introduced CM concept, simulation

- Separate congestion management, protocols
- Let multiple protocols share a single CM
- Separate congestion control and scheduling



# The Congestion Manager

---



- Unmodified receiver network stack
- Single API for congestion management
- Transmissions are orchestrated by the CM

# A simple API

---

API overview:

- open new connections
- request permission to send
- notify of transmission
- update with successes and losses

Plus callbacks:

- send a packet
- rate has changed

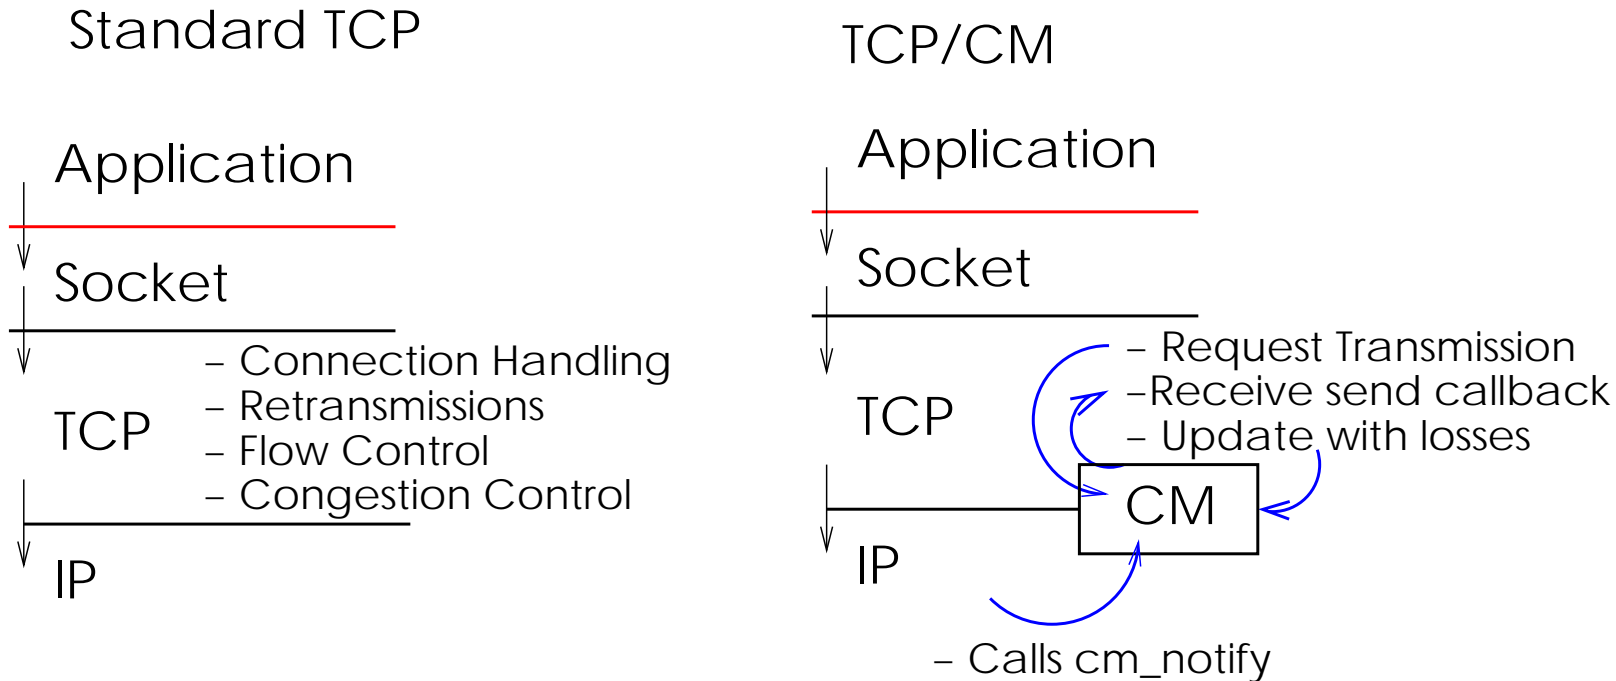
# **Application-controlled transmissions**

✗ Buffering would reduce application control

- “Last-minute adaptation:”
    - TCP (losses)
    - Streaming media (quality)
  - Allows lazy evaluation of work
  - Perfect for in-kernel clients like TCP
- ➔ Used for several very different approaches
- Request/Callback, Rate callbacks, buffered socket

# TCP: Request/Callback API

---

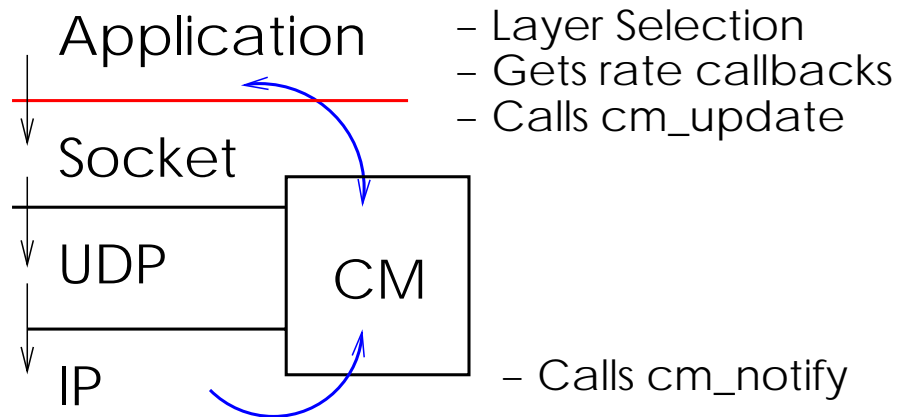


- Gives TCP control over *what* to send

➔ Behavior nearly identical to TCP



# Multi-Rate Encoding Server: **Rate Callbacks**

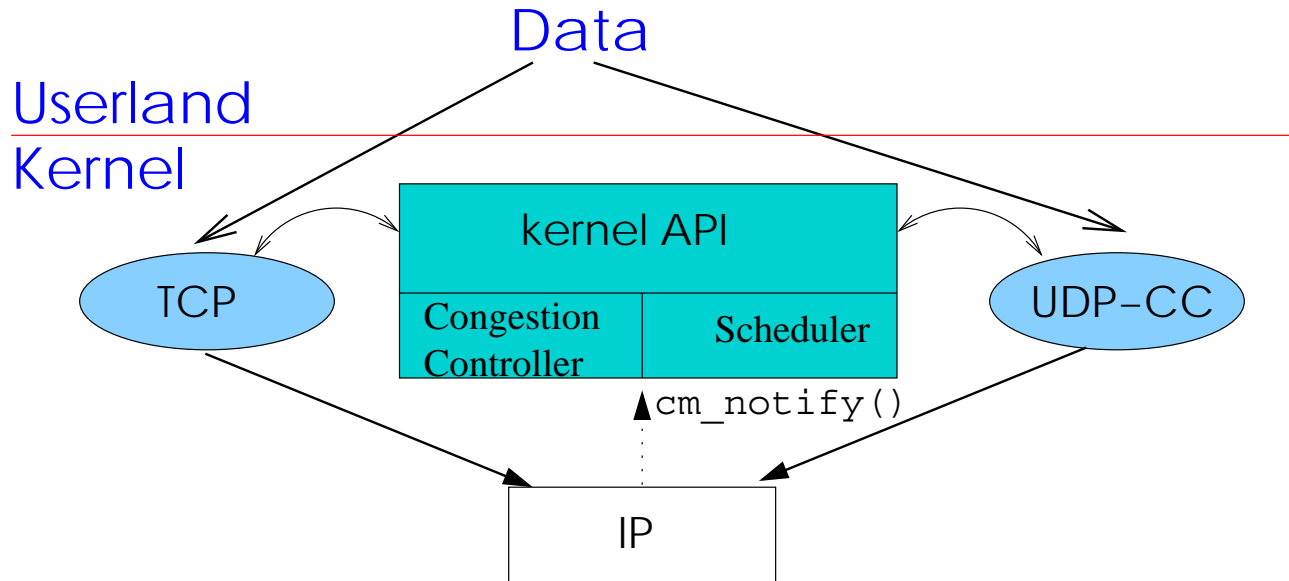


- Select encoding by bandwidth
- Fixed number of encodings more/less bandwidth
- Send at specified rate until notified via callback

→ Reduces extraneous callbacks

# The CM Implementation

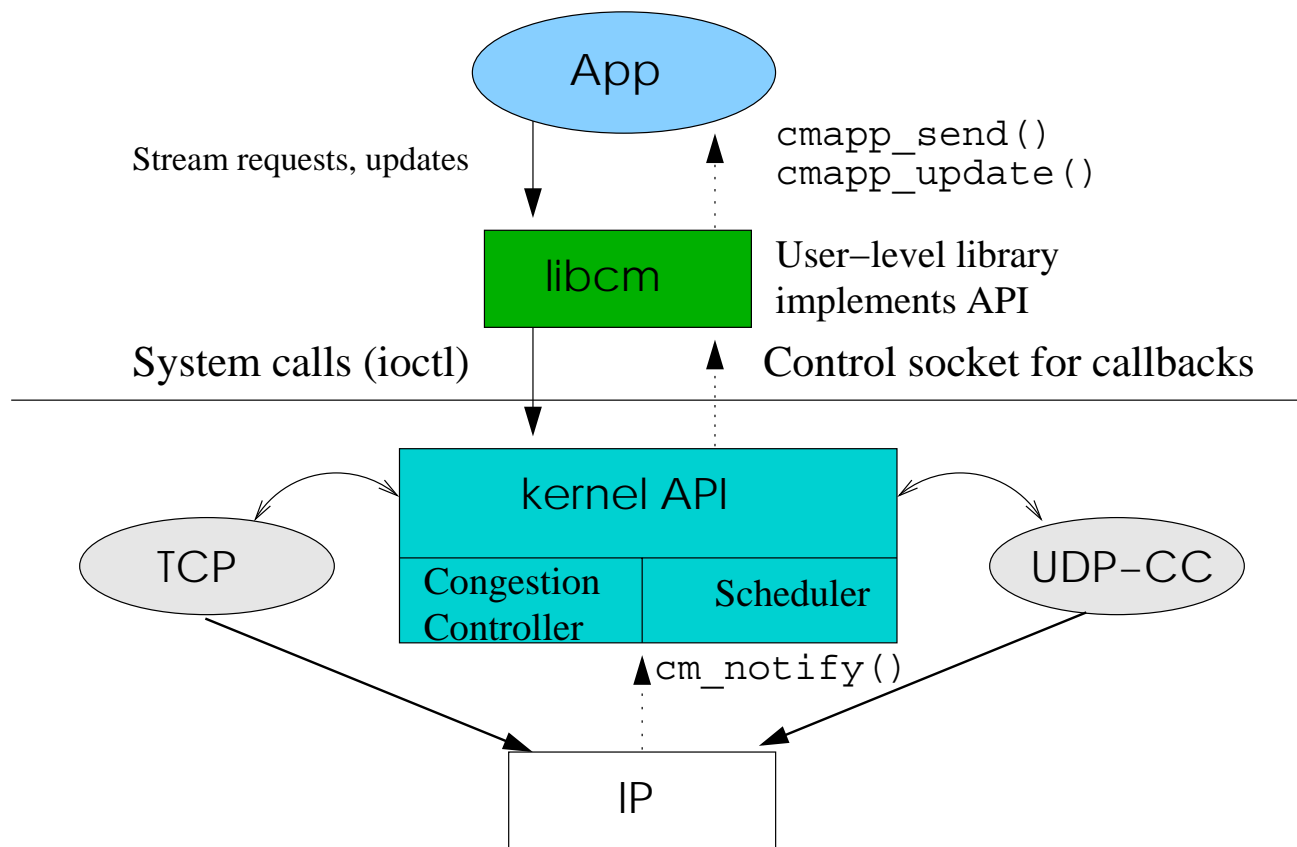
---



- Function callbacks great in kernel
- `ip_output` informs CM of packet transmission using `cm_notify`

# The CM Implementation

---



# Evaluation questions

---

## Impact:

- How does it impact the network?
- How does it impact *my* connections?


## Design & Implementation:

- Can the CM be implemented efficiently?
- Can we write new apps using it?
  - How convenient is the adaptation API?

# The CM is as friendly as TCP

---


Measured throughput vs. loss rate ( $tput \propto \frac{1}{\sqrt{loss}}$ )



- Dummynet random losses
- 10Mbps 60ms RTT
- 10Mbyte file

→ Global network benefits from the CM

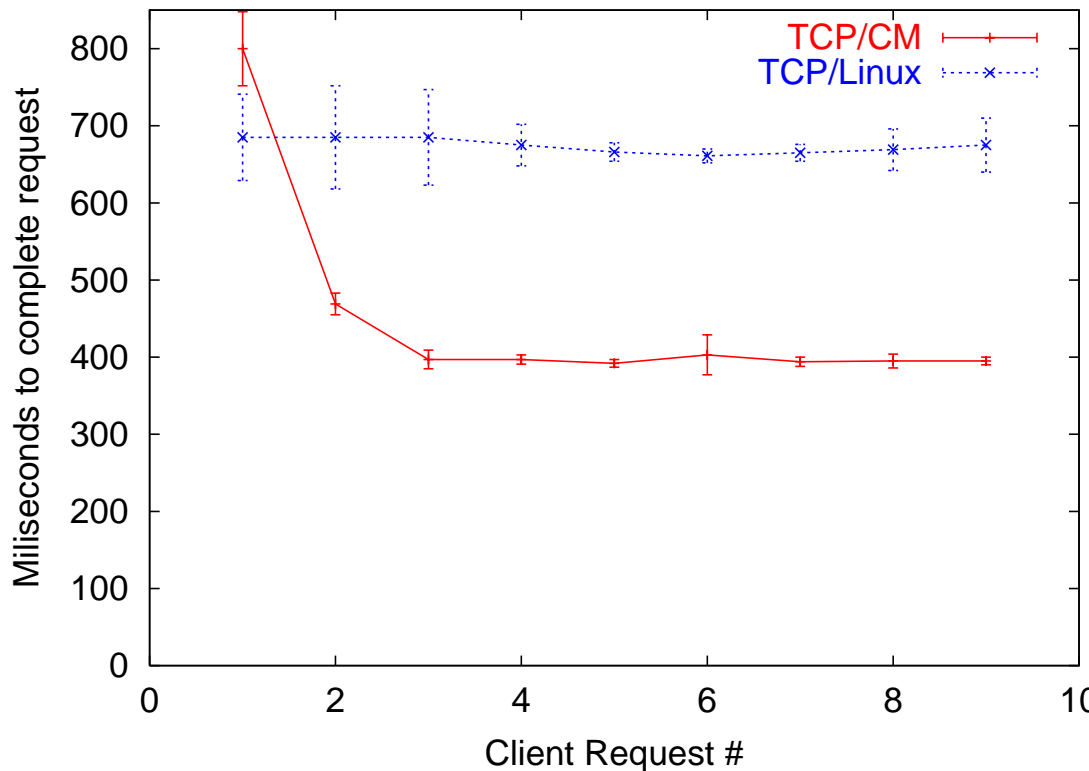
# Testing the effect of flow integration



- Web-like: Issue a request every 500ms  
Regardless of completion of earlier requests
- Measure completion times

# Integrating congestion control helps

Throughput can benefit from sharing

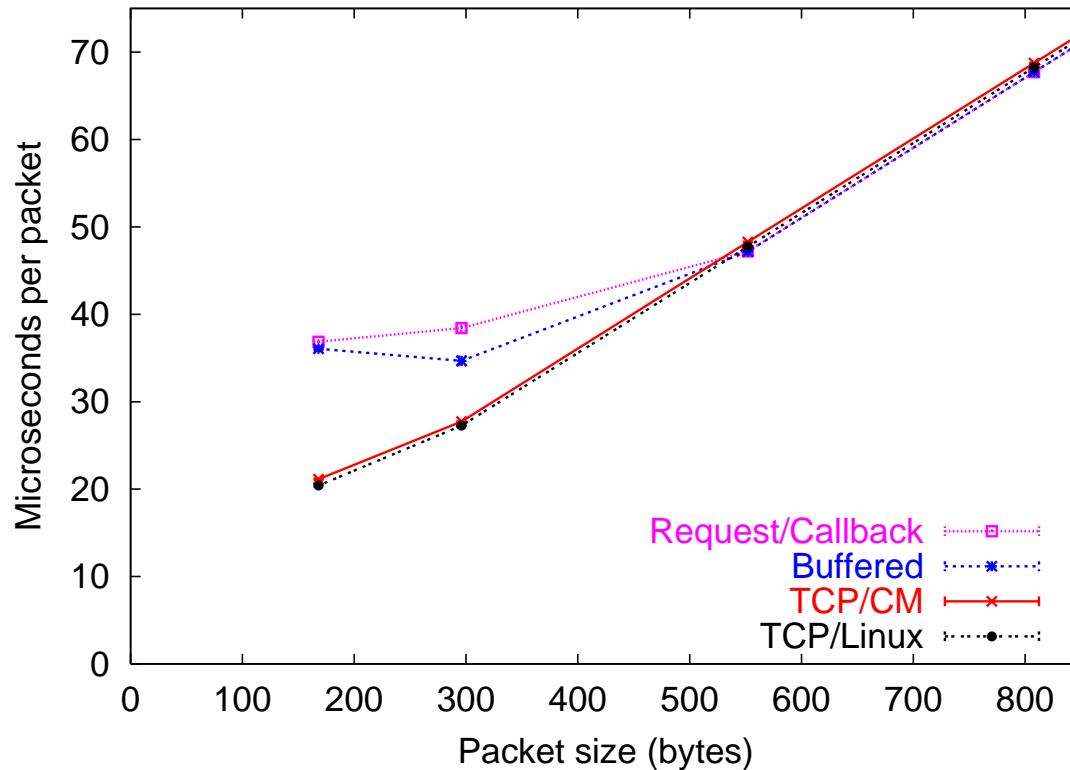


- Web-like workload
- Internet path MIT → Utah
- 500ms request spacing
- 128k file

→ Applications benefit locally from the CM

# The CM is efficient

Examined wall clock time / number of packets



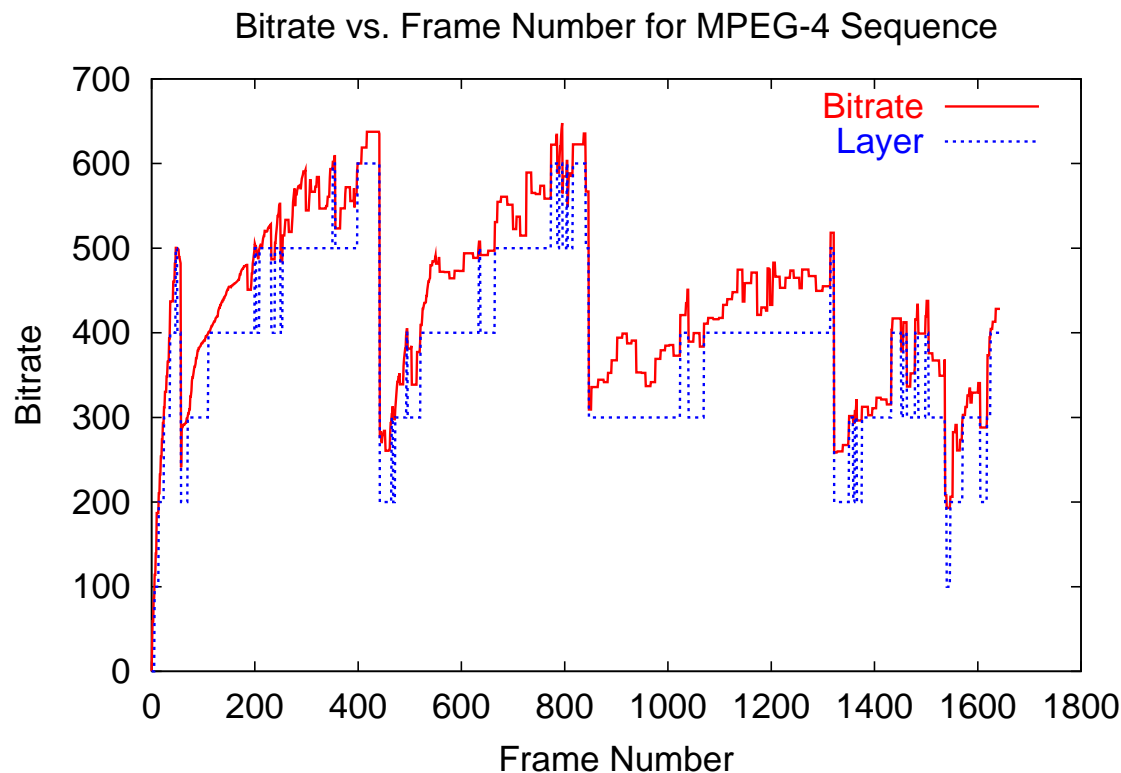
- 100M Ethernet
- 600 Mhz PIIIs
- Quiet network

- < 1-2% difference for in-kernel TCP
- Same throughput with packets > 512 bytes



# The CM enables application to adapt

## Layered MPEG-4 video sender (unrelated project)



- Layer  $n$  at  $100 * n$  Kbps
- Cross traffic from SURGE web workload

→ CM API facilitates adaptive applications

# More applications

---

- Implemented an adaptive vat
  - Audioconferencing tool
  - Avoids wasted packets
  - Achieves same audio quality
- Several adaptive test apps in software release
- ➔ Good platform for research
  - New congestion control algorithms
  - New scheduling algorithms

# **Some Future Work**

---

- Multiuser machine security
- When to aggregate flows
  - More: Between machines
  - Less: QoS flow isolation, proxies
- Feedback frequency and mechanisms

# Conclusions

---

Separating congestion control is *good*

- Potential for global performance gains
- Potential for local performance gains

Our API makes using the CM *easy*

- Add congestion control to non-TCP apps

Our implementation is *efficient*

- Flexibility without loss of efficiency

# **Software and more information**

[nms.lcs.mit.edu/projects/cm/](http://nms.lcs.mit.edu/projects/cm/)

- Freely available (GPL)  
implementation in Linux 2.2
- Internet-Draft of CM spec (ECM WG)
- Lots more!

# **Persistent HTTP**

---

- Protocol-specific solution
- Browsers still open 2-4 streams
- Coupling of fate between unrelated items
- Contributes to protocol complexity

# The CM API

---

- Register a flow  
→ `flow = cm_open(src, dst)`
- Request permission to send  
→ `cm_request(flow)`
- Wait for a callback  
→ `cmapp_send(flow)`
- Transmit up to 1 MTU
- Tell the CM how it went  
→ `cm_update(flow, sent, recd, rtt)`

# The CM API: Requests

---

```
fid = cm_open(struct sockaddr_in *src,  
              struct sockaddr_in *dst)
```

```
mtu = cm_mtu(fid)      Obtain connection's MTU
```

```
cm_request(fid)        Ask to send up to 1 MTU
```

```
cm_register_send(fn)   Set send callback
```

```
cm_register_update(fn) Set update callback
```

```
cm_thresh(down, up)    Set update thresholds
```

```
cm_update(fid, nsent, nrecd, mode, rtt)
```

```
cm_notify(fid, bytes)  Notify # bytes sent
```



## **The CM API: Callbacks**

`cmapp_send(fid)`

→ Application may send 1 MTU on this flow

`cmapp_notify(fid, flow_parameters)`

→ The network conditions for this flow changed